


# What is Caffe?

Convolutional Architecture for Fast Feature Embedding

Caffe is a **deep learning framework** made with expression, speed, and modularity in mind.

It is developed by the **Berkeley Vision and Learning Center (BVLC)** and community contributors. [Yangqing Jia](#) created the project during his PhD at UC Berkeley.



[caffe.berkeleyvision.org](http://caffe.berkeleyvision.org)



[github.com/BVLC/caffe](https://github.com/BVLC/caffe)

<http://caffe.berkeleyvision.org/tutorial/>

## Windows

<http://caffe.berkeleyvision.org/installation.html>

<https://github.com/Microsoft/caffe>

## Linux (Ubuntu)

[http://caffe.berkeleyvision.org/install\\_appt.html](http://caffe.berkeleyvision.org/install_appt.html)



## General Dependencies

## Install Caffe in Linux (ubuntu14.04)

- ☐ `sudo apt-get install libatlas-base-dev`
- ☐ `sudo apt-get install libprotobuf-dev libleveldb-dev libsnappy-dev libopencv-dev libboost-all-dev libhdf5-serial-dev`
- ☐ `sudo apt-get install libgflags-dev libgoogle-glog-dev liblmdb-dev protobuf-compiler`

## Install Caffe

- ☐ `git clone https://github.com/BVLC/caffe.git`
- ☐ `cd caffe`
- ☐ `cp Makefile.config.example Makefile.config`

## CPU-Only (In VMWare)

- ☐ `vi Makefile.config`
- ☐ `# Adjust Makefile.config (for example, if using Anaconda Python)`
- ☐ `# uncomment CPU_ONLY := 1`

## Compile

- ☐ `make all`
- ☐ `make pycaffe`

# Learn LeNet on MNIST

**Dataset** <http://yann.lecun.com/exdb/mnist/>

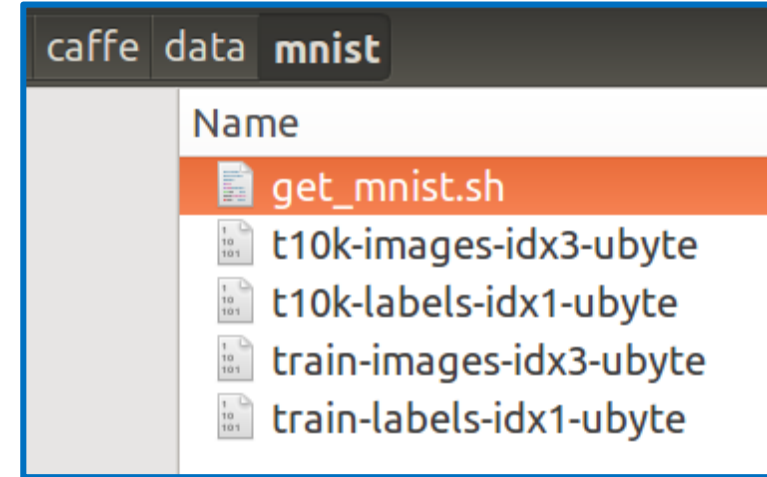

## Training LeNet on MNIST with Caffe

---

We will assume that you have Caffe successfully compiled. If not, please refer to the [Installation page](#). In this tutorial, we will assume that your Caffe installation is located at `CAFFE_ROOT`.

# Test (Mnist)

- ❑ `cd caffe`
- ❑ `sh data/mnist/get_mnist.sh`
- ❑ `sh examples/mnist/create_mnist.sh`
- ❑ `vi examples/mnist/lenet_solver.prototxt`
- ❑ `./examples/mnist/train_lenet.sh`



```
I1124 23:13:49.797040 3455 base_data_layer.cpp:115] Prefetch copied
I1124 23:13:49.798049 3459 data_layer.cpp:102] Prefetch batch: 0 ms.
I1124 23:13:49.798084 3459 data_layer.cpp:103] Read time: 0.12 ms.
I1124 23:13:49.798094 3459 data_layer.cpp:104] Transform time: 0.43 ms.
I1124 23:13:49.953862 3455 base_data_layer.cpp:115] Prefetch copied
I1124 23:13:49.955523 3459 data_layer.cpp:102] Prefetch batch: 1 ms.
I1124 23:13:49.955598 3459 data_layer.cpp:103] Read time: 0.125 ms.
I1124 23:13:49.955606 3459 data_layer.cpp:104] Transform time: 0.749 ms.
I1124 23:13:50.108896 3455 base_data_layer.cpp:115] Prefetch copied
I1124 23:13:50.118820 3459 data_layer.cpp:102] Prefetch batch: 8 ms.
I1124 23:13:50.118875 3459 data_layer.cpp:103] Read time: 0.456 ms.
I1124 23:13:50.118893 3459 data_layer.cpp:104] Transform time: 1.06 ms.
```

```
I1124 23:43:48.829185 3489 base_data_layer.cpp:115] Prefetch copied
I1124 23:43:48.831632 3495 data_layer.cpp:102] Prefetch batch: 1 ms.
I1124 23:43:48.831696 3495 data_layer.cpp:103] Read time: 0.324 ms.
I1124 23:43:48.831702 3495 data_layer.cpp:104] Transform time: 0.935 ms.
I1124 23:43:48.936643 3489 base_data_layer.cpp:115] Prefetch copied
I1124 23:43:48.939280 3495 data_layer.cpp:102] Prefetch batch: 1 ms.
I1124 23:43:48.939393 3495 data_layer.cpp:103] Read time: 0.164 ms.
I1124 23:43:48.939406 3495 data_layer.cpp:104] Transform time: 1.074 ms.
I1124 23:43:49.047644 3489 base_data_layer.cpp:115] Prefetch copied
I1124 23:43:49.049654 3495 data_layer.cpp:102] Prefetch batch: 1 ms.
I1124 23:43:49.049680 3495 data_layer.cpp:103] Read time: 0.161 ms.
I1124 23:43:49.049690 3495 data_layer.cpp:104] Transform time: 0.796 ms.
I1124 23:43:49.153923 3489 solver.cpp:404] Test net output #0: accuracy = 0
.9914
I1124 23:43:49.154065 3489 solver.cpp:404] Test net output #1: loss = 0.029
1876 (* 1 = 0.0291876 loss)
I1124 23:43:49.154078 3489 solver.cpp:322] Optimization Done.
I1124 23:43:49.154080 3489 caffe.cpp:254] Optimization Done.
chgo901@ubuntu:~/caffe$ ^C
```

# Why Caffe? In one sip...

- **Expression:** models + optimizations are plaintext schemas, not code.
- **Speed:** for state-of-the-art models and massive data.
- **Modularity:** to extend to new tasks and settings.
- **Openness:** common code and reference models for reproducibility.
- **Community:** joint discussion and development through BSD-2 licensing.

# Solver.prototxt

<https://github.com/BVLC/caffe/wiki/Solver-Prototxt>

**The solver.prototxt is a configuration file used to tell caffe how you want the network trained.**

- # The train/test net protocol buffer definition
- net: "examples/mnist/lenet\_train\_test.prototxt"
- # test\_iter specifies how many forward passes the test should carry out.
- # In the case of MNIST, we have test batch size 100 and 100 test iterations, covering the full 10,000 testing images.
- test\_iter: 100
- # Carry out testing every 500 training iterations.
- test\_interval: 500
- # The base learning rate, momentum and the weight decay of the network.
- base\_lr: 0.01
- momentum: 0.9
- weight\_decay: 0.0005

```
# The learning rate policy
lr_policy: "inv"
gamma: 0.0001
power: 0.75
# Display every 100 iterations
display: 100
# The maximum number of iterations
max_iter: 10000
# snapshot intermediate results
snapshot: 5000
snapshot_prefix: "examples/mnist/lenet"
# solver mode: CPU or GPU
solver_mode: CPU
```

```

name: "CIFAR10_quick_test"
layer {
  name: "data"
  type: "Input"
  top: "data"
  input_param { shape: { dim: 1 dim: 3
3 dim: 32 dim: 32 } }
}
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
  convolution_param {
    num_output: 32
    pad: 2
    kernel_size: 5
    stride: 1
  }
}

```

```

layer {
  name: "pool1"
  type: "Pooling"
  bottom: "conv1"
  top: "pool1"
  pooling_param {
    pool: MAX
    kernel_size: 3
    stride: 2
  }
}
layer {
  name: "relu1"
  type: "ReLU"
  bottom: "pool1"
  top: "pool1"
}
layer {
  name: "conv2"
  type: "Convolution"
  bottom: "pool1"
  top: "conv2"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
  convolution_param {
    num_output: 32
    pad: 2
    kernel_size: 5
    stride: 1
  }
}

```

```

layer {
  name: "relu2"
  type: "ReLU"
  bottom: "conv2"
  top: "conv2"
}
layer {
  name: "pool2"
  type: "Pooling"
  bottom: "conv2"
  top: "pool2"
  pooling_param {
    pool: AVE
    kernel_size: 3
    stride: 2
  }
}
layer {
  name: "conv3"
  type: "Convolution"
  bottom: "pool2"
  top: "conv3"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
  convolution_param {
    num_output: 64
    pad: 2
    kernel_size: 5
    stride: 1
  }
}

```

```

layer {
  name: "relu3"
  type: "ReLU"
  bottom: "conv3"
  top: "conv3"
}
layer {
  name: "pool3"
  type: "Pooling"
  bottom: "conv3"
  top: "pool3"
  pooling_param {
    pool: AVE
    kernel_size: 3
    stride: 2
  }
}
layer {
  name: "ip1"
  type: "InnerProduct"
  bottom: "pool3"
  top: "ip1"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
  inner_product_param {
    num_output: 64
  }
}

```

```

layer {
  name: "ip2"
  type: "InnerProduct"
  bottom: "ip1"
  top: "ip2"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
  inner_product_param {
    num_output: 10
  }
}
layer {
  name: "prob"
  type: "Softmax"
  bottom: "ip2"
  top: "prob"
}

```

• Train.prototxt

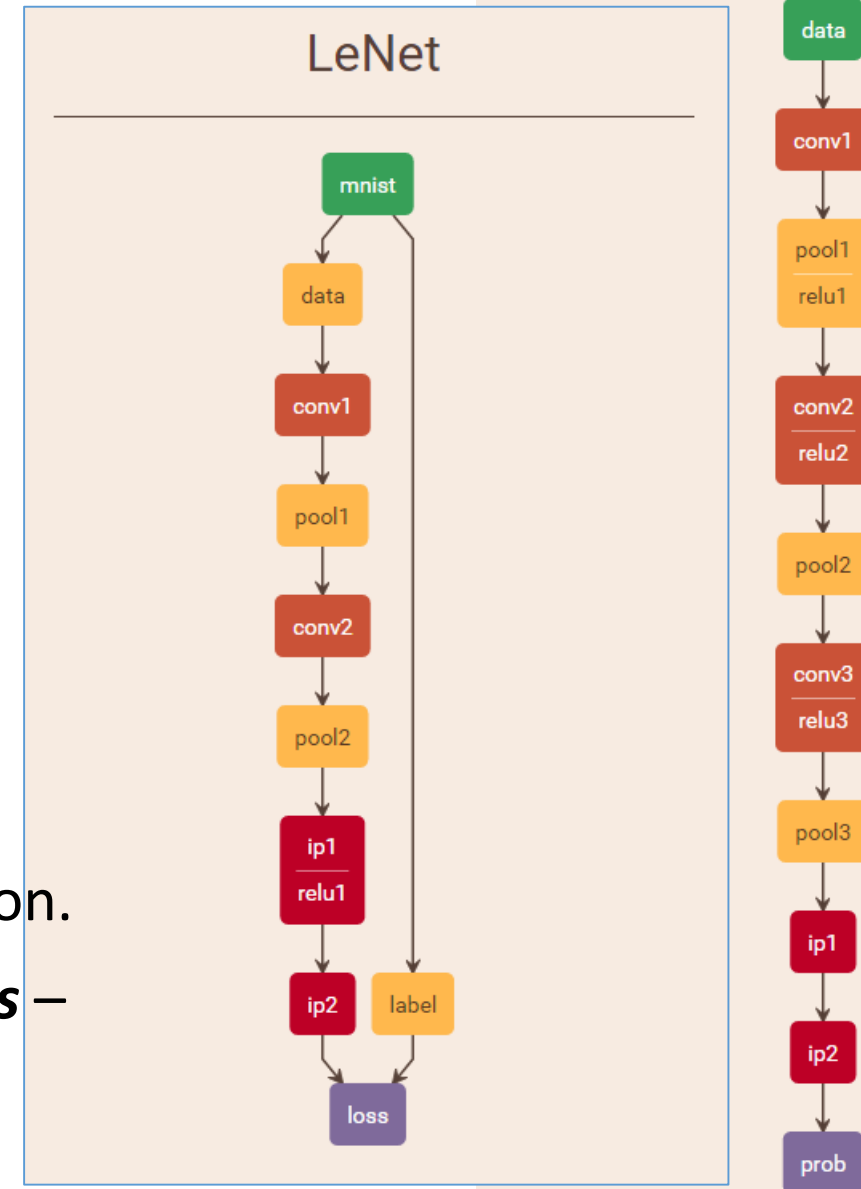
# Net `lenet_train_test.prototxt`

- A network is a set of layers and their connections as a DAG:

```
name: "dummy-net"
layer { name: "data" ...}
layer { name: "conv" ...}
layer { name: "pool" ...}
... more layers ...
layer { name: "loss" ...}
```

- Caffe creates and checks the net from the definition.
- Data and derivatives flow through the net as ***blobs*** – an array interface

## CIFAR10 quick test






# Blob


Blobs are N-D arrays for storing and communicating information.

- hold data, derivatives, and parameters
- lazily allocate memory
- shuttle between CPU and GPU




## Data

Number x  $K$  Channel x *Height* x *Width*  
256 x 3 x 227 x 227 for ImageNet train input



## Parameter: Convolution Weight


$N$  Output x  $K$  Input x *Height* x *Width*  
96 x 3 x 11 x 11 for CaffeNet conv1



## Parameter: Convolution Bias


96 x 1 x 1 x 1 for CaffeNet conv1

```
name: "conv1"  
type: CONVOLUTION  
bottom: "data"  
top: "conv1"  
... definition ...
```



# Blob

Blobs provide a unified memory interface.



**Reshape(num, channel, height, width)**

- declare dimensions
- make *SyncedMem* -- but only lazily allocate

**cpu\_data(), mutable\_cpu\_data()**

- host memory for CPU mode

**gpu\_data(), mutable\_gpu\_data()**

- device memory for GPU mode

**{cpu,gpu}\_diff(), mutable\_{cpu,gpu}\_diff()**

- derivative counterparts to data methods
- easy access to data + diff in forward / backward

# Layer


```
name: "conv1"  
type: CONVOLUTION  
bottom: "data"  
top: "conv1"  
convolution_param {  
  num_output: 20  
  kernel_size: 5  
  stride: 1  
  weight_filler {  
    type: "xavier"  
  }  
}
```

name, type, and the  
connection structure  
(input blobs and  
output blobs)

layer-specific  
parameters

- Every layer type defines

- **Setup**
- **Forward**
- **Backward**



- Nets + Layers are defined by [protobuf](#) schema

# Data Layer

- layer {
  - name: "mnist"
  - type: "Data"
  - transform\_param {
    - scale: 0.00390625 // The input pixels are normalized to [0, 1],  $0.00390625 = 1/256$
  - }
  - data\_param {
    - source: "mnist\_train\_lmdb"
    - backend: LMDB
    - batch\_size: 64 //batch mode
  - }
  - top: "data" //generate two blobs, data blob 和 label blob
  - top: "label"
  - }

# Convolution Layer

[An Explanation of Xavier Initialization](http://andyljones.tumblr.com/post/110998971763/an-explanation-of-xavier-initialization)

<http://andyljones.tumblr.com/post/110998971763/an-explanation-of-xavier-initialization>

- layer {
- name: "conv1"
- type: "Convolution"
- //Parameter adjustment of the learning rate
- param { lr\_mult: 1 }     //the IR of weight is same as the solver
- param { lr\_mult: 2 }     //the LR of bias is twice as high as the solvers'

# Convolution Layer

[An Explanation of Xavier Initialization](http://andyljones.tumblr.com/post/110998971763/an-explanation-of-xavier-initialization)

<http://andyljones.tumblr.com/post/110998971763/an-explanation-of-xavier-initialization>

- convolution\_param {
- num\_output: 20 //number of activation map
- kernel\_size: 5 //the size of filter 5\*5
- stride: 1
- weight\_filler {
- type: "xavier" //xavier algorithm, automatically initialize weight filler,  
//based on the number of input and output neurons
- }
- bias\_filler {
- type: "constant" //initialize bias filler as constant, default value 0
- }
- }
- bottom: "data" //Data Blobs from the underlying layer are used as inputs
- top: "conv1" // produces the `conv1` layer
- }

# Pooling Layer

- layers {
  - name: "pool1"
  - type: POOLING
  - pooling\_param {
    - kernel\_size: 2
    - stride: 2 //Prevent overlapping areas
    - pool: MAX
  - }
  - bottom: "conv1"
  - top: "pool1"}

# Innerproduct Layer (Fully Connection Layer)

- layer {
- name: "ip1" *//caffe calls it an innerproduct layer*
- type: "InnerProduct"
- bottom: "pool3"
- top: "ip1"
- param {
- lr\_mult: 1
- }
- param {
- lr\_mult: 2
- }
- inner\_product\_param {
- num\_output: 64
- }
- }

- layer {
- name: "ip2"
- type: "InnerProduct"
- bottom: "ip1"
- top: "ip2"
- param {
- lr\_mult: 1
- }
- param {
- lr\_mult: 2
- }
- inner\_product\_param {
- num\_output: 10
- }
- }



# ReLU Layer

- layer {
  - name: "relu3"
  - type: "ReLU"
  - bottom: "conv3"
  - top: "conv3"
  - }
- layer {
  - name: "relu2"
  - type: "ReLU"
  - bottom: "conv2"
  - top: "conv2"
  - }
- layer {
  - name: "relu1"
  - type: "ReLU"
  - bottom: "pool1"
  - top: "pool1"
  - }

- Memory Reuse
- The input and output of the blob set to the same name, can be a single element of the operation of the relu save storage space


# Loss layer(Softmax Layer)

- layer {
  - name: "prob"
  - type: "Softmax"
  - bottom: "ip2"
  - top: "prob"
- }

# Loss


<http://caffe.berkeleyvision.org/tutorial/loss.html>

What kind of model is this?




# Loss

What kind of model is this?



Define the task by the **loss**.



Classification  
SoftmaxWithLoss  
HingeLoss

Linear Regression  
EuclideanLoss

Attributes / Multiclassification  
SigmoidCrossEntropyLoss

Others...

New Task  
NewLoss

# The CIFAR-10 dataset

<http://www.cs.toronto.edu/~kriz/cifar.html>

60000 32x32 colour images in 10 classes,  
6000 images per class.  
50000 training images  
10000 test images.

The CIFAR-100 dataset

This dataset is just like the CIFAR-10, except it has 100 classes containing 600 images each. T

**airplane**



**automobile**



**bird**



**cat**



**deer**



**dog**



**frog**



**horse**



**ship**



**truck**



# Share a Sip of Brewed Models

[demo.caffe.berkeleyvision.org](http://demo.caffe.berkeleyvision.org)

demo code open-source and bundled



Maximally accurate	Maximally specific
cat	1.80727
domestic cat	1.74727
feline	1.72787
tabby	0.99133
domestic animal	0.78542

# LAST SIP

Caffe...

- is fast
- is state-of-the-art
- has tips, recipes, demos, and models
- brings together an active community
- ...all for free and open source

Thank you